

CS 7319 Software Architecture

Homework 2: Sentiment Analysis in Two Architectures

Instructor: Dr. Isaac Chow

Fall 2025

Goal: Implement the same functionality using two different software architectures to experience architectural trade-offs directly.

Part A is a single-program design in any language.

Part B is a MapReduce-style pipeline (as in the provided notebook).

1. Learning Outcomes

- Model a problem using two architectural styles and explain the differences.
- Implement a simple keyword-based sentiment analyzer.
- Compare code structure, data flow, complexity, and scalability implications.

2. Problem Statement (common to both parts)

Given a text dataset (one post per line), decide overall mood using keywords:

- Positive keywords: happy, excited, thrilled, love
- Negative keywords: sad, depressed, angry, upset

Classification rules:

- Positive: contains ≥ 1 positive keyword and no negative keyword
- Negative: contains ≥ 1 negative keyword and no positive keyword
- Mixed: contains both positive and negative keywords
- Neutral: otherwise

Data

Provided datasets in the data/ folder:

- data/sample_us_posts.txt (starter)
- data/sample_us_posts_mixed.txt (mixed sentiments)
- data/your_name_posts.txt – (your own posts)

Part A — Single-Program Architecture (language of your choice)

Architectural intent: A simple, single-process program that reads the dataset, scans each line, classifies it, and aggregates totals.

Requirements:

- 1) Choose any language (e.g., Python, Java, C#, ...).

- 2) Read input from a text file (one post per line).
- 3) Tokenize case-insensitively (letters/apostrophes are enough).
- 4) Classify each line using the rules above.
- 5) Aggregate totals per category and print a summary:
 - Counts for Positive/Negative/Mixed/Neutral
 - Verdict: “Happier” (Positive>Negative), “Sadder” (Negative>Positive), or “Tied”
- 6) CLI usage example:
sentiment_prog <path-to-textfile>
- 7) Output example:
Positive=123 Negative=95 Mixed=17 Neutral=64
Verdict: Happier
- 8) (Optional) Produce a simple bar chart.

What to submit for Part A:

1. Screen shots showing that your program is working as intended
2. Brief Description of your architecture, language, and how to build/run
3. Source file(s)

Part B — MapReduce-Style Architecture (as in the existing notebook)

Architectural intent: Model the problem as Map → Shuffle/Group → Reduce to simulate data-parallel processing.

Requirements:

- 1) Map: implement `map(line) → [(label, 1)]` emitting one pair per line (Positive/Negative/Mixed/Neutral).
- 2) Shuffle/Group: group all values by label.
- 3) Reduce: sum values per label to get totals.
- 4) Config knobs (recommended): `USE_COMBINER=True/False`; `NUM_REDUCERS` if you simulate partitions.
- 5) Outputs: same summary and verdict as Part A; if using the notebook, include two charts (counts; Positive vs Negative %).
- 6) Use both datasets (run once per dataset).

What to submit for Part B:

1. Screen shots showing that your program/notebook is working as intended
2. Brief Description of your architecture, explaining where Map, Shuffle, and Reduce occur in your code

3. Completed notebook (or equivalent code)

Comparison & Reflection (Bonus)

Write ½–1 page comparing the two architectures:

- Structure & responsibilities (mapper vs single pass)
- Complexity (which code is simpler? easier to evolve?)
- Scalability & performance (how would each handle 100,000 × data?)

Step-by-Step Instructions

Part A (single program):

- 1) Pick a language you're comfortable with.
- 2) Implement reading → tokenizing → classification → tallying.
- 3) Run on all datasets.
- 4) Print counts and verdict; save output (copy/paste or screenshot).
- 5) (Optional) Generate a bar chart.

Part B (MapReduce-style):

- Using the provided Jupyter Notebook:
 - 1) Launch Jupyter and open `sentiment_mapreduce.ipynb`.
 - 2) (Optional) Use the Dataset Selector to choose a dataset.
 - 3) Run cells top to bottom.
 - 4) Record totals and the verdict.
 - 5) Screenshot the two charts.
- Or implement Map/Shuffle/Reduce in your own language and print the same outputs.